

Predicting gene-disease associations via graph embedding and graph convolutional networks[1]

Lvxing Zhu et. al.

1 Jargon Glossary

DAG - or, *Directed Acyclic Graph*, a graph with each edge having a direction, such that there are NO directed cycles, or closed loops

DNA Sequencing - The general laboratory technique for determining the exact sequence/pattern of nucleotides/bases (Adenine, Cytosine, Guanine, and Thymine) in a DNA molecule

FFNN - Feed Forward Neural Network

GCN - Graph Convolution Network [2]

HMDD - or, *Human MicroRNA-associated Disease Database*, has experimentally supported associations for human micro-RNA and diseases

LDA - or, *Latent Dirichlet Allocation*, given multiple *documents* that contain permutations of *words* in a vocabulary, create topic (representative word) clusters, such that doc. belongs to some topic; an unsupervised, probabilistic model

Pathology - The science of causes and effects of diseases; specifically involves laboratory examination of samples of body tissue

PCA - or, *Principal Components Analysis*, find eigen vectors of the covariance matrix such that highly correlated variable set may be reduced to fewer un-correlated independent variables, in a co-ordinate system whose axes capture all the variance of the original variables; an unsupervised statistical model

Pharmacology - The scientific study of effect of drugs and chemicals on living organisms

Physiology - Physiology is the scientific study of functions and mechanisms in a living system

2 Problem Description

Identifying genes that contribute to a disease can further our understanding in human **physiology**, **pathology** and thus enable better research in **pharmacology**. In this paper, the authors have setup a heterogeneous graph of *genes* and *diseases* and utilised the techniques of graph embedding (DeepWalk + word2vec) followed by GCN [2] to learn embeddings for both the *gene nodes* and the *diseases nodes*. This is followed by a decoder module which learns to predict the probability of association between any *gene-disease* pair. An overview of all the components of the paper is summarised in the *mindmap* in Figure 1.

3 Problem Relevance

Most *in-silico* methods attempt to leverage the *guilt-by-association* principle wherein new gene-disease associations are predicted considering the functional similarity of the genes to known causative genes. Some of the common methods for this prediction are:



Figure 1: Mindmap

1. *Matrix Decomposition*
2. *Network Propagation*
3. *Shallow Machine Learning*
4. *Graph Embedding*

Matrix Decomposition is one of the most common approaches where after decomposing a matrix of known gene-disease association, new associations are recovered by arranging the decompositions for unknown gene-disease pairs. *Network Propagation* relies on PPI graphs and *Shallow ML* methods involve boosted regression trees, again using PPI. *Graph Embedding* approaches are newer, and can incorporate heterogeneous graphs.

A heterogeneous graph can simultaneously contain nodes of type **gene**, **disease**, etc. and the learning task can be setup directly as a *link prediction* problem, between the genes and diseases.

4 Proposed Solution

The authors have setup the downstream task as a link prediction problem, which roughly has 4 stages:

1. Build gene-disease graph
2. Initial node features
3. GCN
4. Decoding

4.1 Building Heterogeneous gene-disease Graph

As depicted in Figure 2 the primary data sources are **DisGeNet** for **gene-disease** linkages, **HumanNet** for **gene-gene** linkages and **MeSH** + **HMDD** for creating the **disease-disease** linkages from the DAG of MeSH. The DAG is used in a process called MISIM [3] that generates scores that can be used for making the disease-only graph.

Once we have the above information, the final heterogeneous graph is constructed as $G = (V, E)$, where $V = V^{gd}$ and $E = E^{gd} \cup E^{gg} \cup E^{dd}$. The adjacency matrix is constructed as follows (Eq. 1, 2, 3 of the paper):

Disclaimer: The equations that follow need reconsideration. I have modified a term.

$$A = \begin{bmatrix} \tilde{A}^{gg} & A^{gd} \\ (A^{gd})^T & \tilde{A}^{dd} \end{bmatrix} \quad (1)$$

where,

$$\begin{aligned} \tilde{A}_{ij}^{gg} &= \phi^i A_{ij}^{gg} \\ \phi^i &= \phi \frac{\sum_j A_{ij}^{gd}}{\sum_j A_{ij}^{gg}} \end{aligned} \quad (2)$$

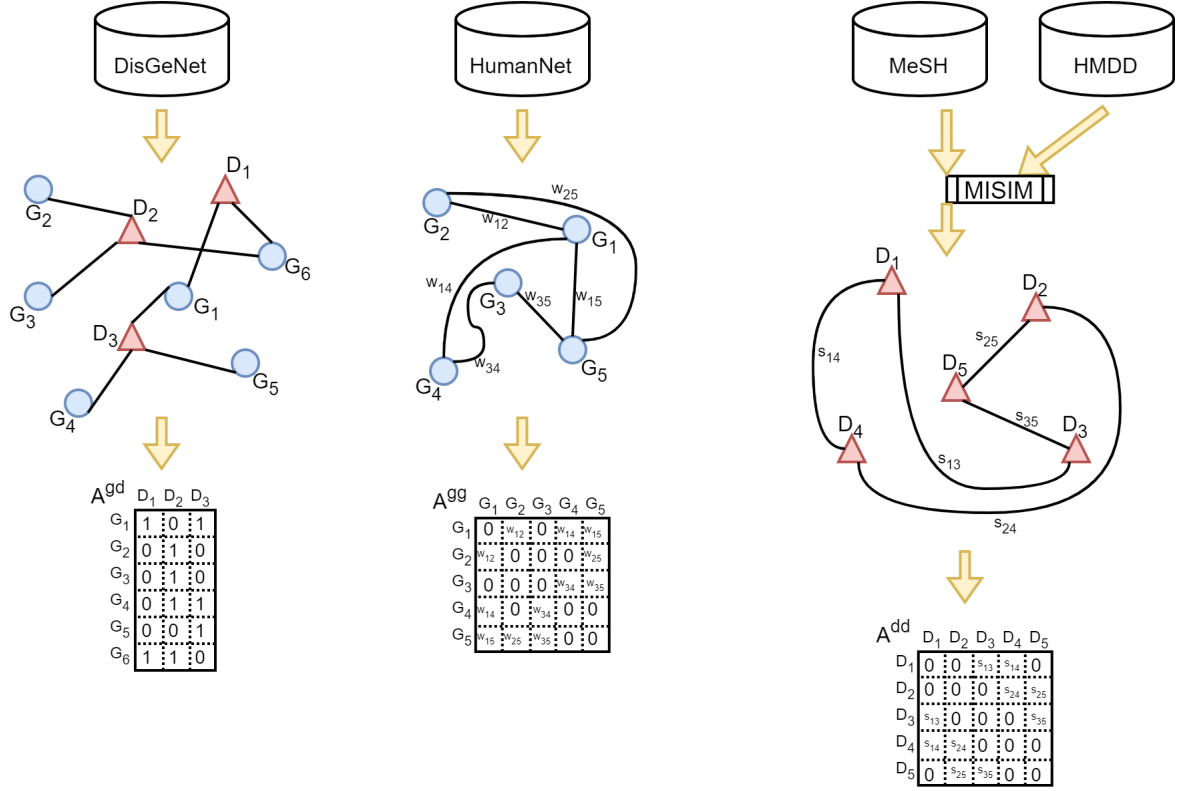


Figure 2: The adjacency matrices for the various edge types. The MISIM [3] method generates the scores for the disease-disease edges. In the final heterogeneous graph structure, the vertex set is retained from the DisGeNet while the edge set is a union of all the above three Adjacencies as per Eq. 1, 2, 3 in the paper.

$$\begin{aligned}\tilde{A}_{ij}^{dd} &= \phi^i A_{ij}^{dd} \\ \phi^i &= \phi \frac{\sum_j A_{ij}^{gd}}{\sum_j A_{ij}^{dd}}\end{aligned}\quad (3)$$

where, ϕ is normalized co-efficients (*authors have not elaborated on the meaning*).

4.2 DeepWalk: generating initial node features

RandomWalk - After the graph is built, each node is assigned an index. For each node, DeepWalk selects λ neighbours based on edge weights from the adjacency matrix A . These λ -length, index sequences form the input to word2vec.

Word2Vec - The index sequences act as documents for a skip-gram version of the word2vec training model. With all the nodes (gene or disease) forming the vocabulary, the index sequences act as sentences. A one-hot vector is assigned to each node (based on its index). Then, for each node, context words are predicted for a fixed window size. The loss function (*Eq. 4 in paper*) compares the prediction against the actual context words in the sequences (*obtained from RandomWalk*)

The above two steps generate the initial feature vectors for each node in the heterogeneous graph. The point to note is that these features only capture the **structural** information of the graph.

The GCN network by itself can not learn anything, unless there is a down- stream task. My understanding is that the decoder layer is essentially the downstream task for the GCN.

4.3 Node embeddings via Graph Convolution Network

At this point, we have a graph (heterogeneous) and initial node-features. Kipf's GCN [2] has a good application in this scenario to learn even better connectivity embeddings for the nodes. Through message passing via the *symmetric degree normalized* adjacency matrix, each node eventually learns about all other nodes in the graph (over multiple iterations, all hop information is spread). All that is needed is a downstream task.

$$embeddings = \sigma(\tilde{D}^{\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} X \Theta) \quad (4)$$

4.4 Decoding novel gene-disease pairs

For the downstream task, the authors concatenate embeddings of a gene and a disease and let a 3-layer FFNN predict the probabilities for a valid association or an invalid association (essentially a binary classification). For this, the set of positive connections and negative connections are drawn from the graph G . The loss function is a combination of **negative log-likelihood** and a novel **cluster loss** (*which looks at approximate clusters based on cosine similarity of the embedding vectors*). The clusters in question are approximate in the sense that a node is assigned to a cluster based on a *cluster centre* calculated within a distance threshold.

Gene Prioritization - It essentially means assigning ranks to genes based on weighted associations to a particular disease. The first step is to list all genes with **valid** associations (from section 4.4) for some disease. Next, filter the list based on a threshold on the association probability score. The final step is to sort the filtered gene list according to their DeepWalk probability score. The result is the final output (*Table III in the paper is an example*).

5 Positive Points

- One of the most straight-forward application of Graph Neural Networks for the task of gene prioritization
- For the most part, the paper is well described

6 Negative Points

- The construction of the combined adjacency matrix (Eq. 1 in the paper) follows some logic, however, it is not revealed in the paper
- The candidate gene list is selected from the decoding module, but, the sorting is done based on the DeepWalk scores. Authors do not provide a justification.
- The authors have mentioned that they used **dropout on Adjacency matrix** as a means to overcome over-fitting of the GCN. This approach doesn't seem right as random edge dropping breaks the domain context.

7 Questions

1. The cluster loss uses a threshold; apparently it is based on distance but, how is it selected?
2. How is the threshold during gene prioritization selected?

References

- [1] L. Zhu, Z. Hong, and H. Zheng, "Predicting gene-disease associations via graph embedding and graph convolutional networks," in *2019 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*, pp. 382–389, 2019.
- [2] T. N. Kipf and M. Welling, "Semi-supervised classification with graph convolutional networks," in *International Conference on Learning Representations*, 2017.
- [3] D. Wang, J. Wang, M. Lu, F. Song, and Q. Cui, "Inferring the human microRNA functional similarity and functional network based on microRNA-associated diseases," *Bioinformatics*, vol. 26, pp. 1644–1650, 05 2010.